



Nearest Neighbor Methods

Rishabh Iyer

University of Texas at Dallas

Based on the slides of Vibhav Gogate, Nick Rouzzi, David Sontag and few other sources

Nearest Neighbor Methods




- Learning
 - Store all training examples

$$\mathcal{T} = \underbrace{\{(x_1, y_1), \dots, (x_M, y_M)\}}_{M \text{ Train Data points}}$$


- Classifying a new point x'
 - Find the training example $(x^{(i)}, y^{(i)})$ such that $x^{(i)}$ is closest (for some notion of close) to x'
 - Classify x' with the label $y^{(i)}$

$x' \rightarrow$ Test point.


Nearest Neighbor Methods




Nearest Neighbor Methods



Nearest Neighbor Methods



Nearest Neighbor Methods




Nearest Neighbor Methods



1-NN : Closest.


K-NN : k Closest points

$x^1 : \{y^1, \dots, y^k\}$




k -nearest neighbor methods look at the k closest points in the training set and take a majority vote
(should choose k to be odd)

Nearest Neighbor Methods

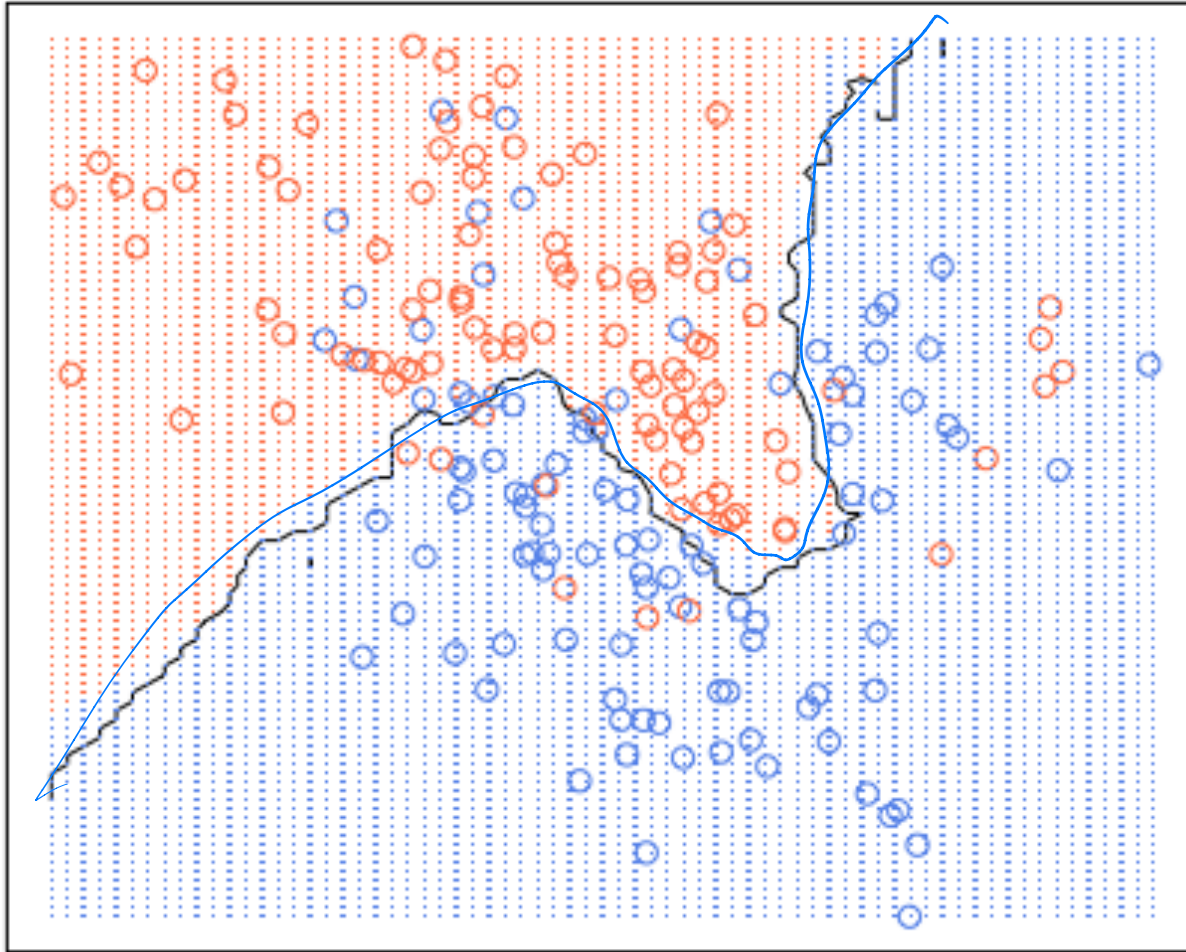


k -nearest neighbor methods look at the k closest points in the training set and take a majority vote (should choose k to be odd)

1-NN Example


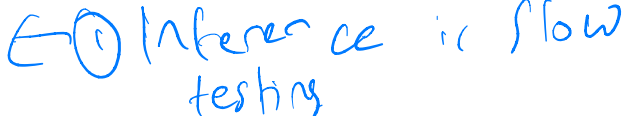



20-NN Example



Nearest Neighbor Methods



- Applies to data sets with points in \mathbb{R}^d
 - Best for large data sets with only a few (< 20) attributes
- Advantages
 - Learning is easy  *no training*
 - Can learn complicated decision boundaries
- Disadvantages
 - Classification is slow (need to keep the entire training set around) $M \sim |M|$  *Inference is slow testing*
 - Easily fooled by irrelevant attributes  *2 High memory*

- How to choose the right measure of closeness?
 - Euclidean distance is popular, but many other possibilities

- How to pick k ?
 - Too small and the estimates are noisy, too large and the accuracy suffers

$$x', x^{(i)} \in \mathbb{R}^d$$

- What if the nearest neighbor is really far away?

$$\begin{array}{ccc} \parallel x' - x^{(i)} \parallel_2^2 & = & \sum_{j=1}^d (x'_j - x_j^{(i)})^2 \\ \uparrow & & \uparrow \\ \text{test} & & \text{train} \end{array}$$

Choosing the Distance



- Euclidean distance makes sense when each of the features is roughly on the same scale
 - If the features are very different (e.g., ⁵height and ²⁰⁻³⁰age), then Euclidean distance makes less sense as height would be less significant than age simply because age has a larger range of possible values
 - To correct for this, feature vectors are often recentered around their means and scaled by the standard deviation over the training set

Normalization

[Mean - Variance Normalization]



- Sample mean

$$\bar{x} = \frac{1}{n} \sum_{i=1}^n x^{(i)}$$

- Sample variance (biased)

$$\hat{\sigma}_k^2 = \frac{1}{n} \sum_{i=1}^n (x_k^{(i)} - \bar{x}_k)^2$$

x_k
 \downarrow
 $x_k = \frac{x_k - \bar{x}_k}{\sigma_k}$

Min-Max Normalization


$$x_k^{\max} = \max_{i=1:n} x_k$$

$$x_k^{\min} = \min_{i=1:n} x_k \quad (i)$$


$$x_k'' = \frac{x_k - x_k^{\min}}{x_k^{\max} - x_k^{\min}} \in [0, 1]$$

Same Normalization for train
& test

$X \leftarrow \text{test}$



$$\mu_{tr}^1 = \frac{\sum_{tr} y^{train}}{G^{train}}$$




$$\mu_{te}^1 = \frac{\sum_{te} y^{train}}{G^{train}}$$

Irrelevant Attributes




Consider the nearest neighbor problem in one dimension



Irrelevant Attributes





Now, add a new attribute that is just random noise...



$$1 \rightarrow 2 \rightarrow 3$$

$$d_{12} \leq d_{13}$$



K-Dimensional Trees



inference / testing

- In order to do classification, we can compute the distances between all points in the training set and the point we are trying to classify
- With ^{1M} m data points in ¹⁰⁰ n -dimensional space, this takes $O(mn)$ time for Euclidean distance
- It is possible to do better if we do some preprocessing on the training data


Test: n_t ~ 100M (slow?)
Prediction ~ MS

- k-d trees provide a data structure that can help simplify the classification task by constructing a tree that partitions the search space
 - Starting with the entire training set, choose some dimension, i
 - Select an element of the training data whose i^{th} dimension has the median value among all elements of the training set
 - Divide the training set into two pieces: depending on whether their i^{th} attribute is smaller or larger than the median
 - Repeat this partitioning process on each of the two new pieces separately

K-Dimensional Trees





- Building a K-D tree from training data:
 - $\{(1,9), (2,3), (4,1), (3,7), (5,4), (6,8), (7,2), (8,8), (7,9), (9,6)\}$
 - pick random dimension, find median, split data, repeat



K-Dimensional Trees



- Building a K-D tree from training data:
 - $\{(1,9), (2,3), (4,1), (3,7), (5,4), (6,8), (7,2), (8,8), (7,9), (9,6)\}$
 - pick random dimension, find median, split data, repeat





K-Dimensional Trees: Inference









- Find NNs for new point (7,4) \leftarrow Test
 - find region containing (7,4)
 - compare to all points in region

Approximate NN



- By design, the constructed k-d tree is “bushy”
 - The idea is that if new points to classify are evenly distributed throughout the space, then the expected (amortized) cost of classification is approximately $O(d \log n)$ operations
- Summary
 - k-NN is fast and easy to implement
 - No training required
 - Can be good in practice (where applicable)

K-mediod.



$$10 = m.$$

Color: {red, green, blue, ... }

→ [1 0 0 ... 0]

1-Hot Encoding.

[0 0 1 0 0 ... 0]

