Neural Networks

Rishabh Iyer

University of Texas at Dallas

Thanks to Nick Rouzzi for sharing these slides

Handwritten Digit Recognition

- Given a collection of handwritten digits and their corresponding labels, we'd like to be able to correctly classify handwritten digits

  - A simple algorithmic technique can solve this problem with 95% accuracy

    - State-of-the-art methods can achieve near 99% accuracy (you've probably seen these in action if you've deposited a check recently)
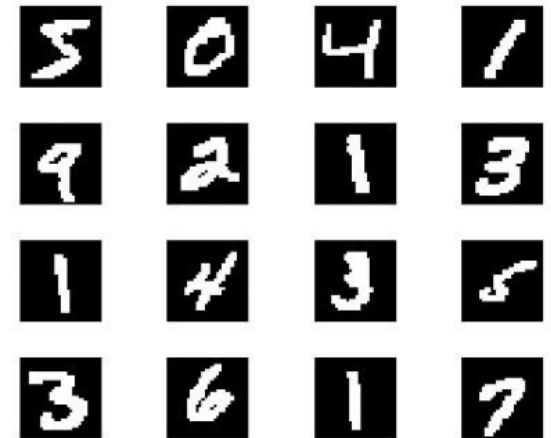

Digits from the MNIST data set

Neural Networks

- The basis of neural networks was developed in the 1940s -1960s

  - The idea was to build mathematical models that might "compute" in the same way that neurons in the brain do

  - As a result, neural networks are biologically inspired, though many of the algorithms developed for them are not biologically plausible

  - Perform surprisingly well for the handwritten digit recognition task (and many others)

Neural Networks

- Neural networks consist of a collection of artificial neurons

- There are different types of neuron models that are commonly studied

  - The perceptron (one of the first studied)

  - The sigmoid neuron (one of the most common, but many more)

  - Rectified linear units

- A neural network is a directed graph consisting of a collection of neurons (the nodes), directed edges (each with an associated weight), and a collection of fixed binary inputs
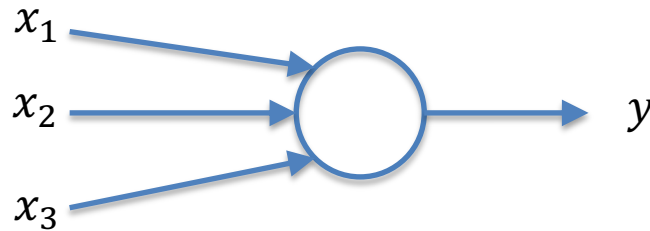
The Perceptron

- A perceptron is an artificial neuron that takes a collection of binary inputs and produces a binary output

    - The output of the perceptron is determined by summing up the weighted inputs and thresholding the result: if the weighted sum is larger than the threshold, the output is one (and zero otherwise)
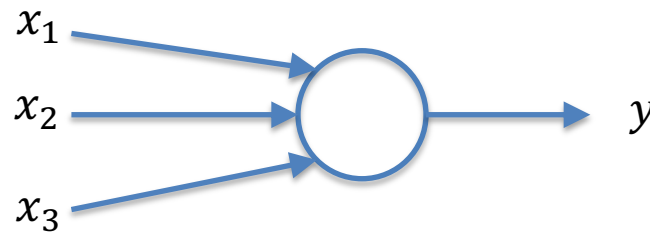
$$y = \begin{cases} 1 & w_1 x_1 + w_2 x_2 + w_3 x_3 > threshold \\ 0 & otherwise \end{cases}$$
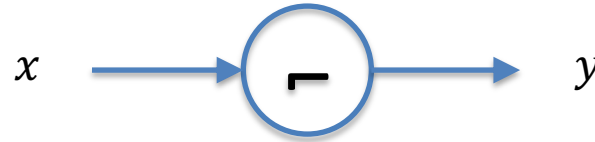
Perceptrons

- Perceptrons are usually expressed in terms of a collection of input weights and a bias b (which is the negative threshold)

$$y = \begin{cases} 1 & w_1 x_1 + w_2 x_2 + w_3 x_3 + b > 0 \\ 0 & otherwise \end{cases}$$

- A single node perceptron is just a linear classifier

- This is actually where the "perceptron algorithm" comes from
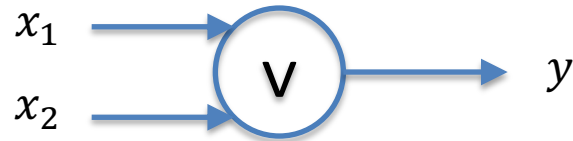
Perceptron for NOT

$$x \longrightarrow \boxed{ \llcorner } \longrightarrow y$$

- Choose $w = -1$, threshold $= -.5$

- $y = \begin{cases} 1 & -x > -.5 \\ 0 & -x \leq -.5 \end{cases}$
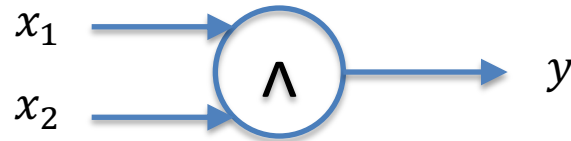
Perceptron for OR

Perceptron for OR

- Choose $w_1 = w_2 = 1$, threshold $= 0$

- $y = \begin{cases} 1 & x_1 + x_2 > 0 \\ 0 & x_1 + x_2 \leq 0 \end{cases}$

Perceptron for AND

Perceptron for AND

$x_1 \longrightarrow$

$x_2 \longrightarrow$ $\wedge$ $\longrightarrow y$

- Choose $w_1 = w_2 = 1$, threshold $= 1.5$

- $y = \begin{cases} 1 & x_1 + x_2 > 1.5 \\ 0 & x_1 + x_2 \leq 1.5 \end{cases}$
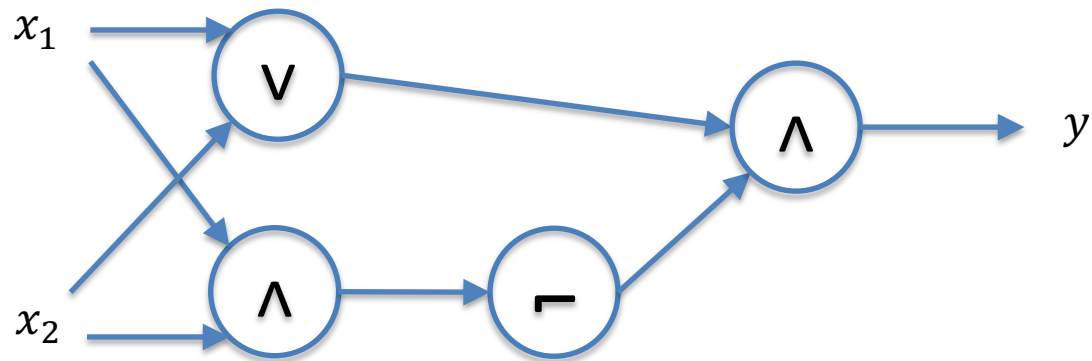
Perceptron for XOR

Perceptron for XOR

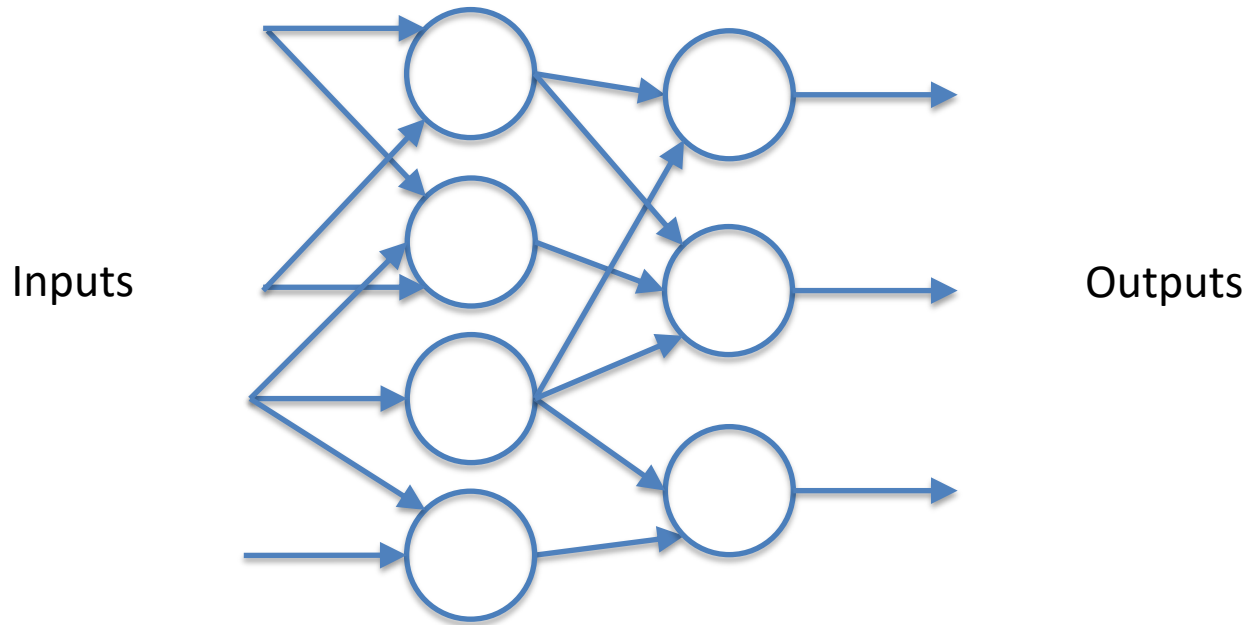- Need more than one perceptron!

- Weights for incoming edges are chosen as before

- Networks of perceptrons can encode any circuit!

Neural Networks

- Gluing a bunch of perceptrons together gives us a neural network

- In general, neural nets have a collection of binary inputs and a collection of binary outputs

Inputs

Outputs

Beyond Perceptrons

- Given a collection of input-output pairs, we'd like to learn the weights of the neural network so that we can correctly predict the output of an unseen input

  - We could try learning via gradient descent (e.g., by minimizing the Hamming loss)

    - This approach doesn't work so well: small changes in the weights can cause dramatic changes in the output

    - This is a consequence of the discontinuity of sharp thresholding (same problem we saw with perceptron alg.)

The Sigmoid Neuron

- A sigmoid neuron is an artificial neuron that takes a collection of real inputs and produces an output in the interval [0,1]

  - The output is determined by summing up the weighted inputs plus the bias and applying the sigmoid function to the result
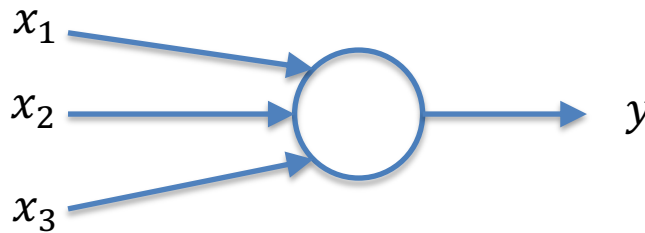
$$y = \sigma(w_1 x_1 + w_2 x_2 + w_3 x_3 + b)$$

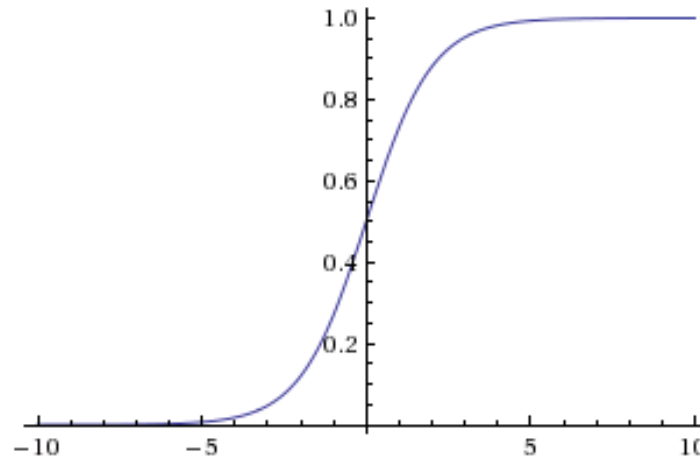where σ is the sigmoid function

The Sigmoid Function

- The sigmoid function is a continuous function that approximates a step function

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

Rectified Linear Units

- The sigmoid neuron approximates a step function as a smooth function

- The relu is given by $\max(0, x)$ which can be approximated as a smooth continuous function $\ln(1 + e^x)$
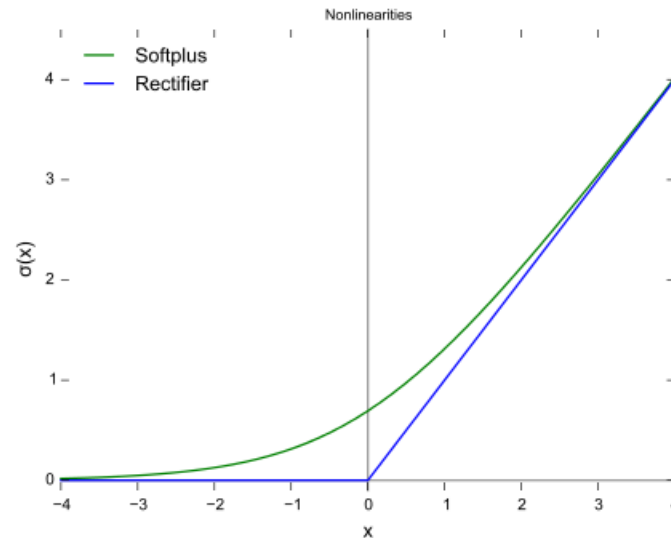
Softmax

- The softmax function maps a vector of real numbers to a vector of probabilities as

$$\text{softmax}(z)_j = \frac{e^{z_j}}{\sum_k e^{z_k}}$$

- If there is a dominant value in z, then it will become one under the softmax

- Often used as the final layer of a neural network

Multilayer Neural Networks
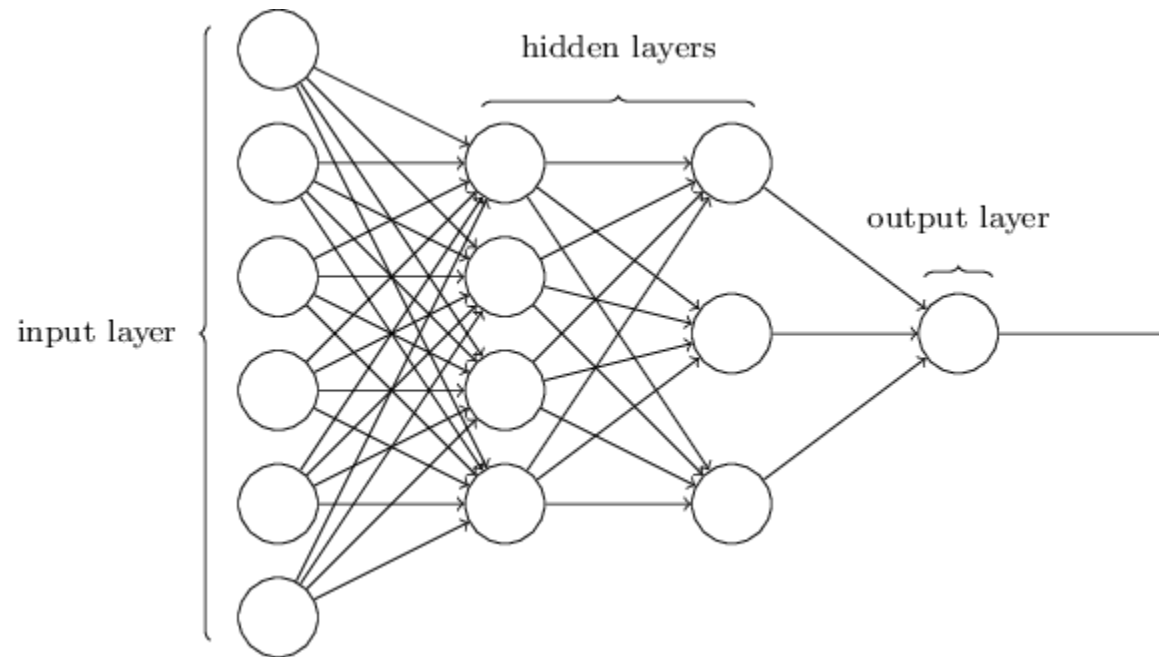
from Neural Networks and Deep Learning by Michael Nielson

Multilayer Neural Networks

NO intralayer connections
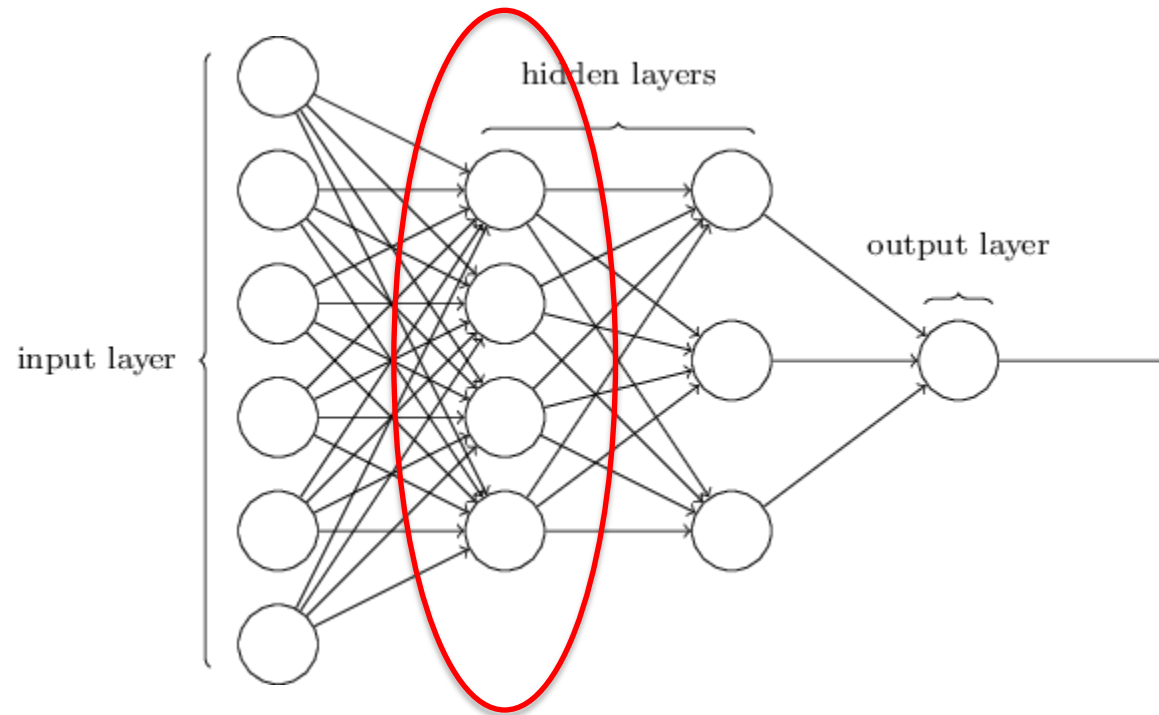
from Neural Networks and Deep Learning by Michael Nielson

Neural Network for Digit Classification

from Neural Networks and Deep Learning by Michael Nielson

Neural Network for Digit Classification

from Neural Networks and Deep Learning by Michael Nielson

Expressiveness of NNs
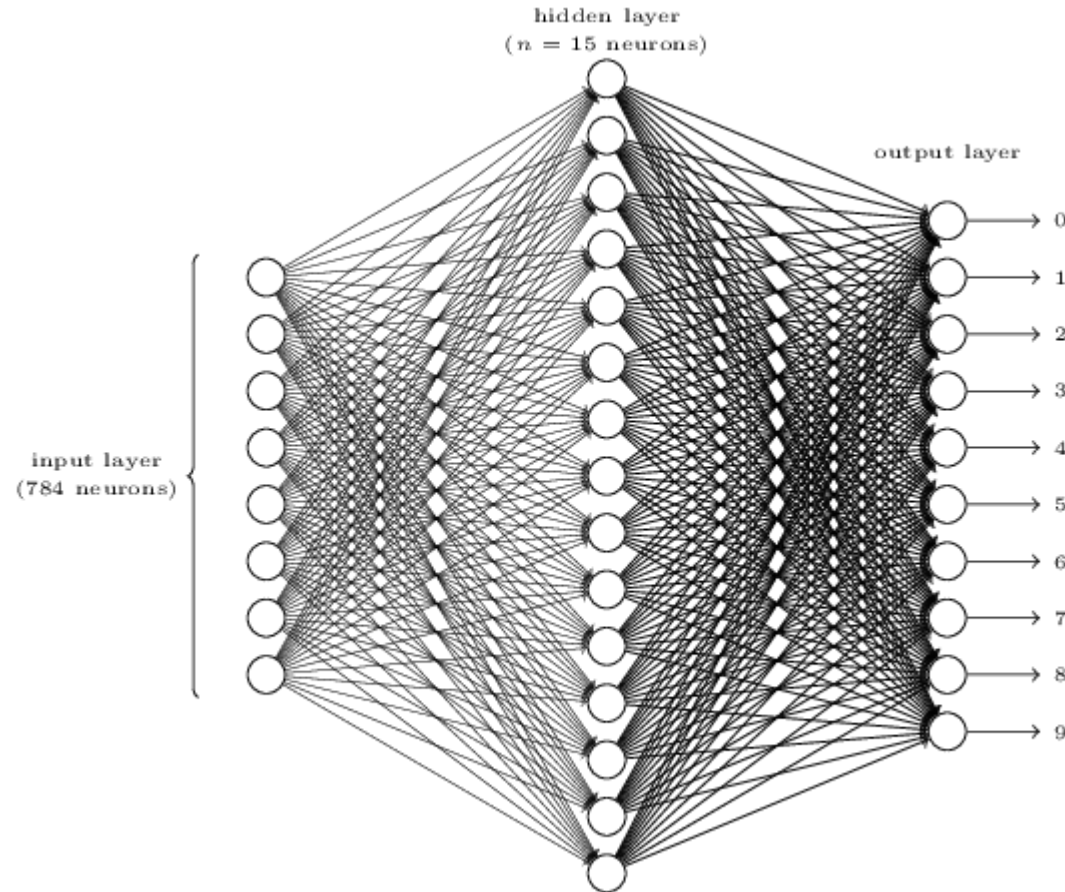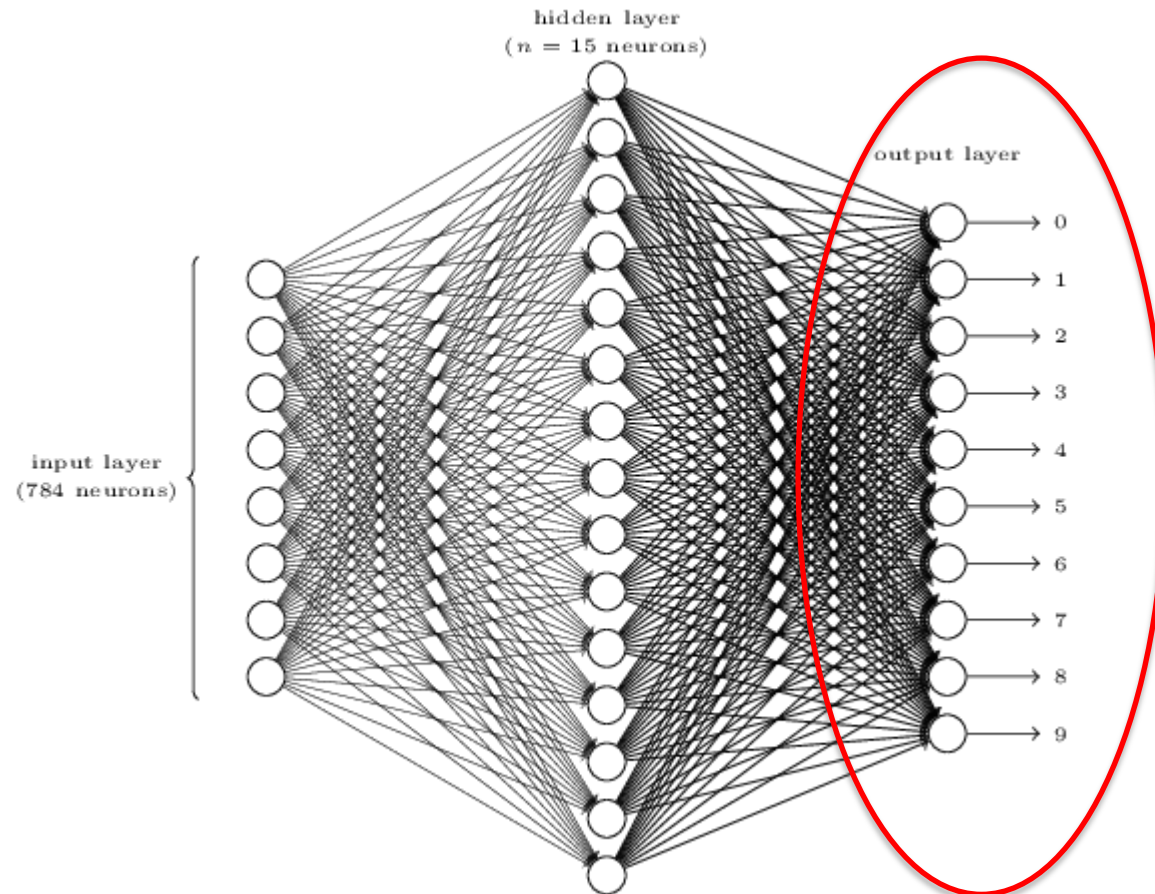
- Boolean functions

  - Every Boolean function can be represented by a network with a single hidden layer consisting of possibly exponentially many hidden units

- Continuous functions

  - Every bounded continuous function can be approximated up to arbitrarily small error by a network with one hidden layer

  - Any function can be approximated to arbitrary accuracy with two hidden layers

Expressiveness of NNs

- **Theorem [Zhang et al. 2016]:** There exists a two-layer neural network with ReLU activations and $2n + d$ weights that can represent any function on a sample of size n in d dimensions

  - This should mean that it is very easy to overfit with neural networks

  - Generalization performance of networks is difficult to assess theoretically

Training Neural Networks

- To do the learning, we first need to define a loss function to minimize

$$C(w, b) = \frac{1}{2M} \sum_m \| y^m - a(x^m, w, b) \|^2$$

- The training data consists of input output pairs $(x^1, y^1), \dots, (x^M, y^M)$

- $a(x^m, w, b)$ is the output of the neural network for the m^{th} sample

- w and b are the weights and biases

Gradient of the Loss

- The derivative of the loss function is calculated as follows

$$\frac{\partial C(w, b)}{\partial w_k} = \frac{1}{M} \sum_m [y^m - a(x^m, w, b)] \frac{\partial a(x^m, w, b)}{\partial w_k}$$

  - To compute the derivative of a, use the chain rule and the derivative of the sigmoid function

$$\frac{d\sigma(z)}{dz} = \sigma(z) \cdot (1 - \sigma(z))$$

  - This gets complicated quickly with lots of layers of neurons

Stochastic Gradient Descent

- To make the training more practical, stochastic gradient descent is used instead of standard gradient descent

- Recall, the idea of stochastic gradient descent is to approximate the gradient of a sum by sampling a few indices and averaging

$$\nabla_x \sum_{i=1}^{n} f_i(x) \approx \frac{1}{K} \sum_{k=1}^{K} \nabla_x f_{i^k}(x)$$

here, for example, each i^k is sampled uniformly at random from $\{1, \ldots, n\}$

Computing the Gradient

- We'll compute the gradient for a single sample

$$C(w, b) = \frac{1}{2} \|y - a(x, w, b)\|^2$$

- Some definitions:

  - L is the number of layers

  - a_j^l is the output of the j^{th} neuron on the l^{th} layer

  - z_j^l is the weighted input of the j^{th} neuron on the l^{th} layer

$$z_j^l = \sum_k w_{jk}^l a_k^{l-1} + b_j^l$$

  - δ_j^l is defined to be $\frac{\partial C}{\partial z_j^l}$

Computing the Gradient

For the output layer, we have the following partial derivative

$$\frac{\partial C}{\partial z_j^L} = -\left(y_j - a_j^L\right)\frac{\partial a_j^L}{\partial z_j^L}$$

$$= -\left(y_j - a_j^L\right)\frac{\partial \sigma\left(z_j^L\right)}{\partial z_j^L}$$

$$= -\left(y_j - a_j^L\right)\sigma\left(z_j^L\right)\left(1 - \sigma\left(z_j^L\right)\right)$$

$$= \delta_j^L$$

- For simplicity, we will denote the vector of all such partials for each node in the l^{th} layer as δ^l

Computing the Gradient

For the $L - 1$ layer, we have the following partial derivative

$$\frac{\partial C}{\partial z_k^{L-1}} = \sum_j (a_j^L - y_j) \frac{\partial a_j^L}{\partial z_k^{L-1}}$$

$$= \sum_j (a_j^L - y_j) \frac{\partial \sigma(z_j^L)}{\partial z_k^{L-1}}$$

$$= \sum_j (a_j^L - y_j) \, \sigma(z_j^L) \left(1 - \sigma(z_j^L)\right) \frac{\partial z_j^L}{\partial z_k^{L-1}}$$

$$= \sum_j (a_j^L - y_j) \, \sigma(z_j^L) \left(1 - \sigma(z_j^L)\right) \frac{\partial \sum_{k'} w_{jk'}^L a_{k'}^{L-1} + b_j^L}{\partial z_k^{L-1}}$$

$$= \sum_j (a_j^L - y_j) \, \sigma(z_j^L) \left(1 - \sigma(z_j^L)\right) \sigma(z_k^{L-1}) \left(1 - \sigma(z_k^{L-1})\right) w_{jk}^L$$

$$= \left((\delta^L)^T w_{*k}^L\right) \left(1 - \sigma(z_k^{L-1})\right) \sigma(z_k^{L-1})$$

Computing the Gradient

- We can think of w^l as a matrix

- This allows us to write

$$\delta^{L-1} = \left((\delta^L)^T w^L\right)^T \circ \left(1 - \sigma(z^{L-1})\right) \circ \sigma(z^{L-1})$$

  where $\sigma(z^{L-1})$ is the vector whose k^{th} component is $\sigma\left(z_k^{L-1}\right)$

- Applying the same strategy, for $l < L$

$$\delta^l = \left(\left(\delta^{l+1}\right)^T w^{l+1}\right)^T \circ \left(1 - \sigma(z^l)\right) \circ \sigma(z^l)$$

Computing the Gradient

- Now, for the partial derivatives that we care about

$$\frac{\partial C}{\partial b_j^l} = \frac{\partial C}{\partial z_j^l} \cdot \frac{\partial z_j^l}{\partial b_j^l} = \delta_j^l$$

$$\frac{\partial C}{\partial w_{jk}^l} = \frac{\partial C}{\partial z_j^l} \cdot \frac{\partial z_j^l}{\partial w_{jk}^l} = \delta_j^l a_k^{l-1}$$

- We can compute these derivatives one layer at a time!

Backpropagation

- Compute the inputs/outputs for each layer by starting at the input layer and applying the sigmoid functions

- Compute δ^L for the output layer

$$\delta_j^L = -\left(y_j - a_j^L\right)\sigma\left(z_j^L\right)\left(1 - \sigma\left(z_j^L\right)\right)$$

- Starting from $l = L - 1$ and working backwards, compute

$$\delta^l = \left(\left(\delta^{l+1}\right)^T w^{l+1}\right)^T \circ \sigma\left(z^l\right) \circ \left(1 - \sigma\left(z^l\right)\right)$$
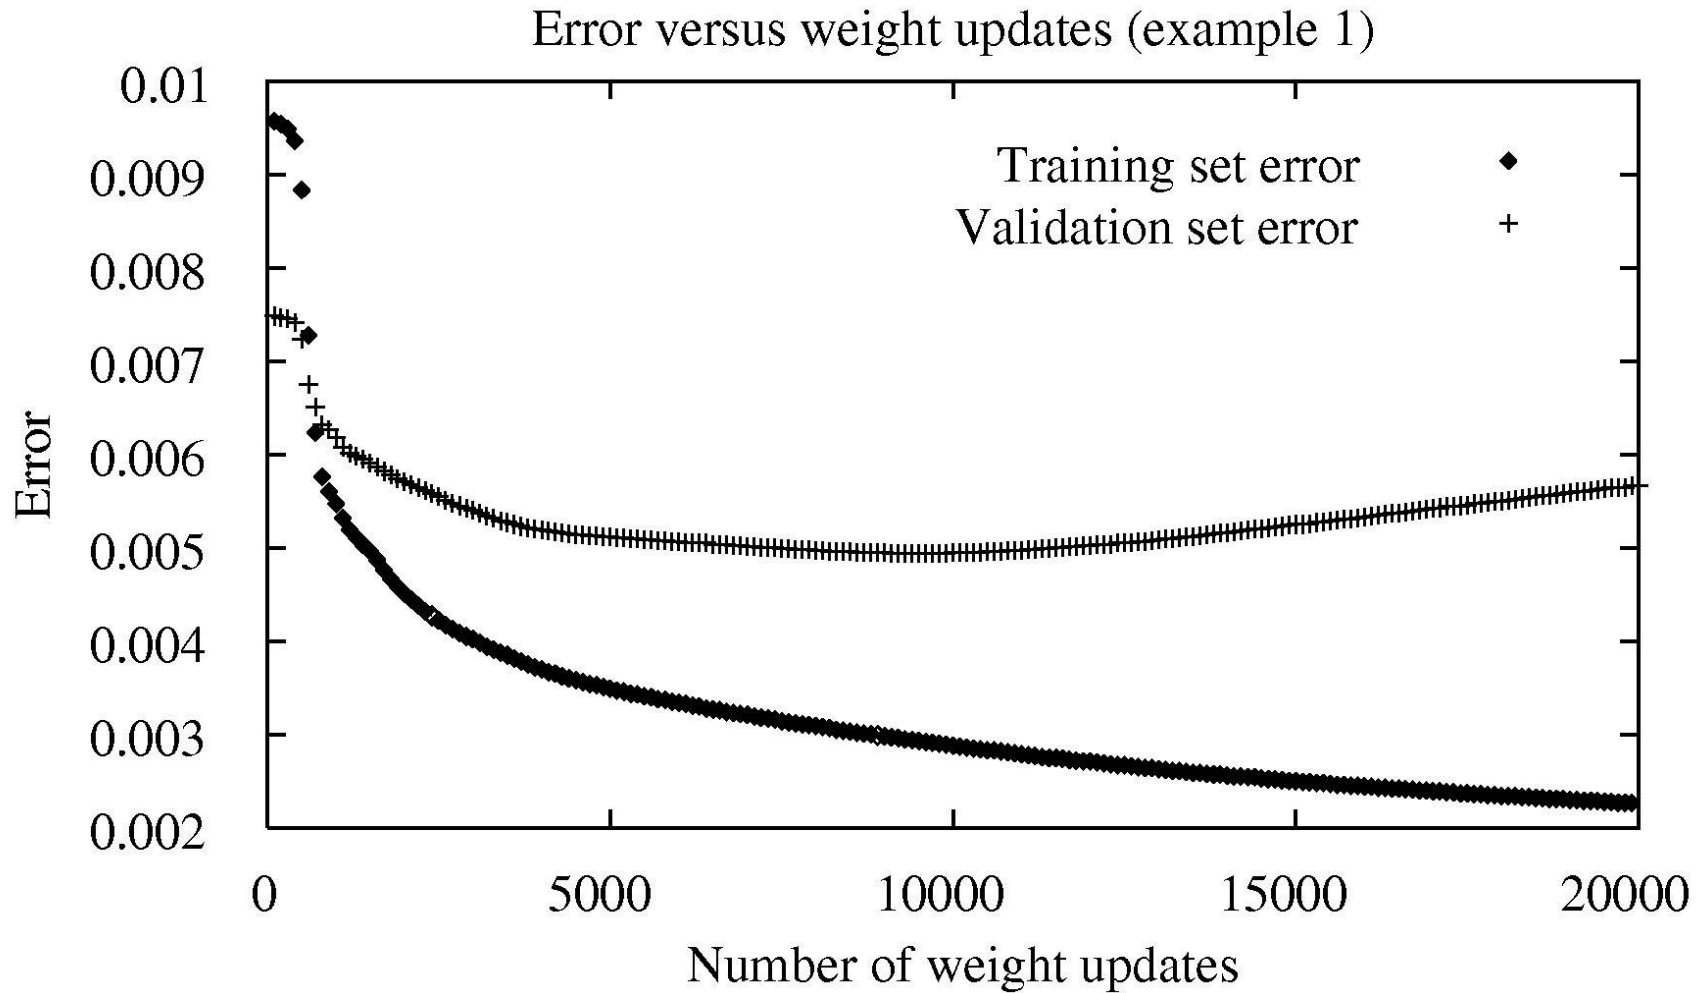
- Perform gradient descent

$$b_j^l = b_j^l - \gamma \cdot \delta_j^l$$

$$w_{jk}^l = w_{jk}^l - \gamma \cdot \delta_j^l a_k^{l-1}$$

Backpropagation

- Backpropagation converges to a local minimum (loss is not convex in the weights and biases)

  - Like EM, can just run it several times with different initializations

  - Training can take a very long time (even with stochastic gradient descent)

  - Prediction after learning is fast

  - Sometimes include a momentum term α in the gradient update

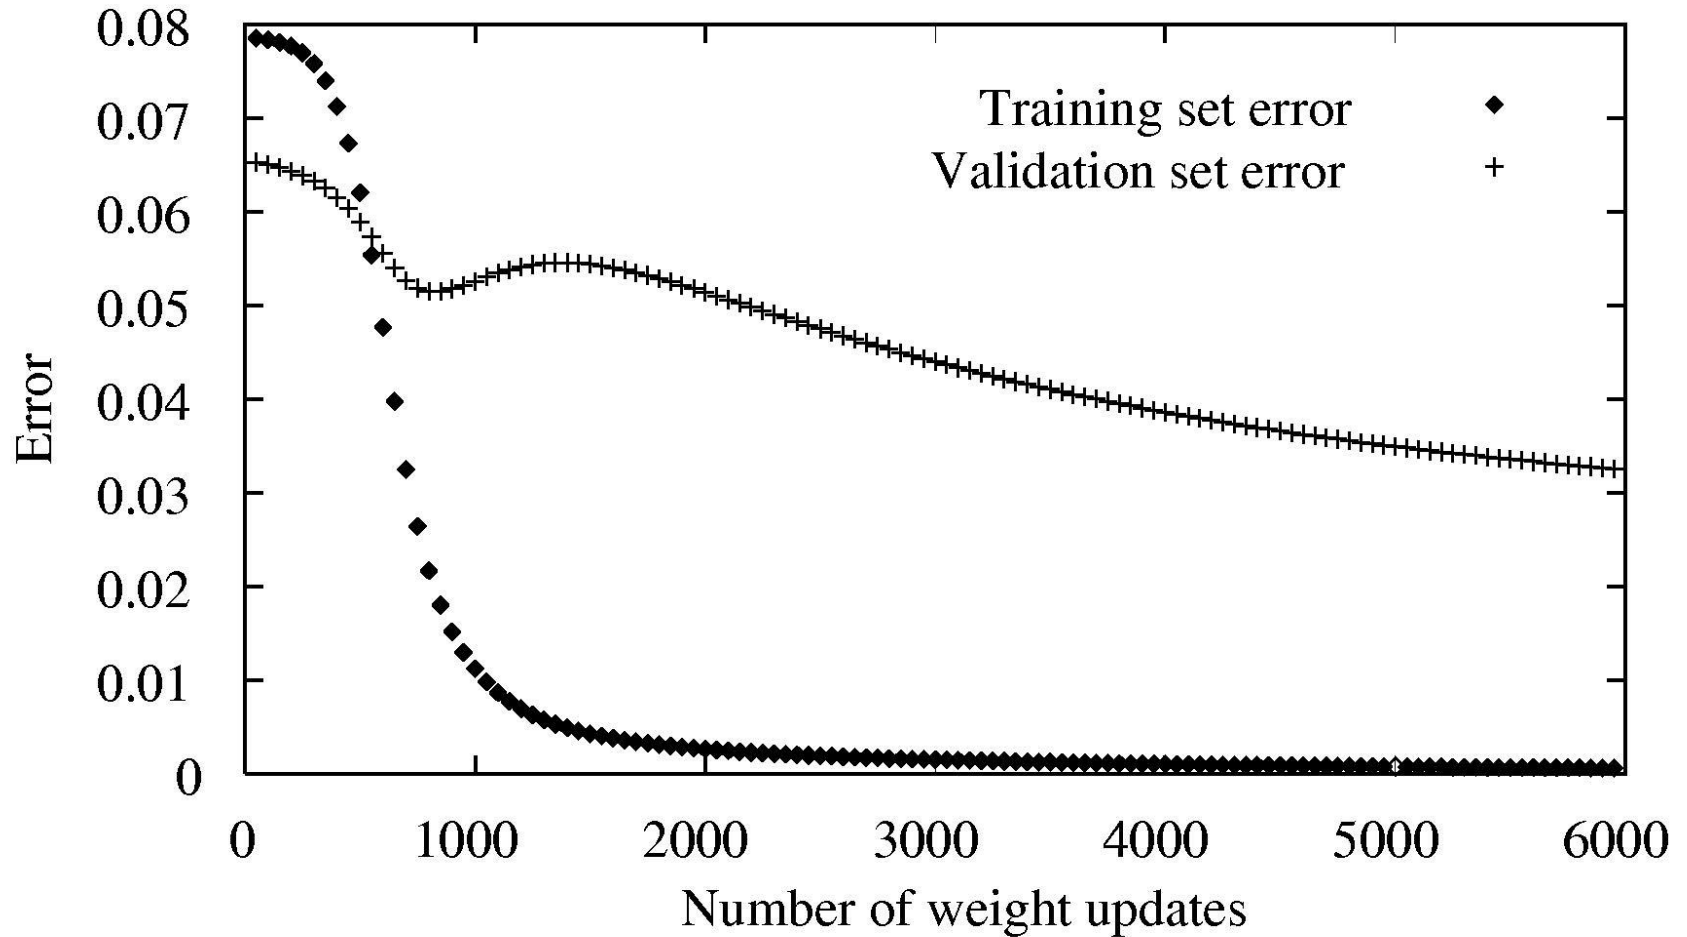$$w(t) = w(t-1) - \gamma \cdot \nabla_w C(t-1) + \alpha(-\gamma \cdot \nabla_w C(t-2))$$

Overfitting

Error versus weight updates (example 1)

Overfitting

Error versus weight updates (example 2)

Neural Networks in Practice

- Many ways to improve weight learning in NNs

  - Use a regularizer! (better generalization?)

  - Try other loss functions, e.g., the cross entropy

    $$-y \log a(x, w, b) - (1 - y) \log(1 - a(x, w, b))$$

  - Initialize the weights of the network more cleverly

    - Random initializations are likely to be far from optimal

- The learning procedure can have numerical difficulties if there are a large number of layers

Regularized Loss

- Penalize learning large weights

$$C'^{(w,b)} = \frac{1}{2M} \sum_m \|y^m - a(x^m, w, b)\|^2 + \frac{\lambda}{2} \|w\|_2^2$$

- Can still use the backpropagation algorithm in this setting

- ℓ_1 regularization can also be useful

- Regularization can help with convergence, λ should be chosen with a validation set

Dropout

- A heuristic bagging-style approach applied to neural networks to counteract overfitting

  - Randomly remove a certain percentage of the neurons from the network and then train only on the remaining neurons

  - The networks are recombined using an approximate averaging technique (keeping around too many networks and doing proper bagging can be costly in practice)

Other Techniques

- Early stopping

  - Stop the learning early in the hopes that this prevents overfitting

- Parameter tying

  - Assume some of the weights in the model are the same to reduce the dimensionality of the learning problem

  - Also a way to learn "simpler" models

  - Can lead to significant compression in neural networks (i.e., >90%)

Other Ideas

- Convolutional neural networks

  - Instead of the output of every neuron at layer l being used as an input to every neuron at layer $l + 1$, the edges between layers are chosen more locally

  - Many tied weights and biases (i.e., convolution nets apply the same process to many different local chunks of neurons)

  - Often combined with pooling layers (i.e., layers that, say, half the number of neurons by replacing small regions of neurons with their maximum output)

  - Used extensively for image classification tasks