



# Lecture 5

## SVMs with Slack

### (Not Linearly Separable)

Rishabh Iyer

University of Texas at Dallas

# Recap SVMs



$$\min_{w,b} \|w\|^2 \leftarrow \frac{1}{\|w\|} \text{ (margin)}$$


Max margin

such that

$$y^{(i)}(w^T x^{(i)} + b) \geq 1, \text{ for all } i$$

- This is a standard quadratic programming problem
  - Falls into the class of **convex optimization problems**
  - Can be solved with many specialized optimization tools (e.g., `quadprog()` in MATLAB)

# Recap SVMs



- Where does the name come from?
  - The set of all data points such that  $y^{(i)}(w^T x^{(i)} + b) = 1$  are called **support vectors**
  - The SVM classifier is completely determined by the support vectors (you could delete the rest of the data and get the same answer)

# Dual SVM = Original SVM Formulation



$$\max_{\lambda \geq 0} -\frac{1}{2} \sum_i \sum_j \lambda_i \lambda_j y_i y_j \underbrace{(x^{(i)})^T (x^{(j)})}_{\text{inner product}} + \sum_i \lambda_i$$

such that

$$\sum_i \lambda_i y_i = 0$$

←  $\lambda^*$   
→  $(w^*, b^*)$

- This is the same as original SVM formulation!
- The dual formulation only depends on inner products between the data points

$$\max_{\lambda \geq 0} -\frac{1}{2} \sum_i \sum_j \lambda_i \lambda_j y_i y_j \underbrace{\Phi(x^{(i)})^T}_{\text{inner product}} \underbrace{\Phi(x^{(j)})}_{\text{inner product}} + \sum_i \lambda_i$$

such that


$$\sum_i \lambda_i y_i = 0$$

- The dual formulation only depends on inner products between the data points
  - Same thing is true if we use feature vectors instead

Projected 4D

---

$y = f(x)$



$$\begin{aligned} \min \quad & f(x) \\ \text{s.t.} \quad & x \geq x_1 \\ & x \leq x_2 \end{aligned}$$

# The Kernel Trick



- For some feature vectors, we can compute the inner products quickly, even if the feature vectors are very large
- This is best illustrated by example

- Let  $\phi(\underline{x_1, x_2}) = \begin{bmatrix} x_1 x_2 \\ x_2 x_1 \\ x_1^2 \\ x_2^2 \end{bmatrix}$

- $$\begin{aligned} \phi(x_1, x_2)^T \phi(z_1, z_2) &= x_1^2 z_1^2 + 2x_1 x_2 z_1 z_2 + x_2^2 z_2^2 \\ &= (x_1 z_1 + x_2 z_2)^2 \\ &= (x^T z)^2 \end{aligned}$$

# The Kernel Trick



- For some feature vectors, we can compute the inner products quickly, even if the feature vectors are very large
- This is best illustrated by example

- Let  $\phi(x_1, x_2) = \begin{bmatrix} x_1 x_2 \\ x_2 x_1 \\ x_1^2 \\ x_2^2 \end{bmatrix}$

- $\underbrace{\phi(x_1, x_2)^T \phi(z_1, z_2)} = x_1^2 z_1^2 + 2x_1 x_2 z_1 z_2 + x_2^2 z_2^2$   
 $= (x_1 z_1 + x_2 z_2)^2$   
 $= (x^T z)^2$

Reduces to a dot product in the original space



# The Kernel Trick



- The same idea can be applied for the feature vector  $\phi$  of all polynomials of degree (exactly)  $d$   
 $k(x, y) = (x^T y)^d$ 
  - $\phi(x)^T \phi(z) = (x^T z)^d$
- More generally, a **kernel** is a function  $k(x, z) = \phi(x)^T \phi(z)$  for some feature map  $\phi$
- Rewrite the dual objective

$$\max_{\lambda \geq 0, \sum_i \lambda_i y_i = 0} -\frac{1}{2} \sum_i \sum_j \lambda_i \lambda_j y_i y_j \boxed{k(x^{(i)}, x^{(j)})} + \sum_i \lambda_i$$

# Examples of Kernels



- Polynomial kernel of degree exactly  $d$ 
  - $k(x, z) = (x^T z)^d$
- General polynomial kernel of degree  $d$  for some  $c$ 
  - $k(x, z) = (x^T z + c)^d$
- Gaussian kernel for some  $\sigma$ 
  - $k(x, z) = \exp\left(\frac{-\|x-z\|^2}{2\sigma^2}\right)$
  - The corresponding  $\phi$  is infinite dimensional!
- So many more...

- Consider the Gaussian kernel

$$\begin{aligned}\exp\left(\frac{-\|x - z\|^2}{2\sigma^2}\right) &= \exp\left(\frac{-(x - z)^T(x - z)}{2\sigma^2}\right) \\ &= \exp\left(\frac{-\|x\|^2 + 2x^T z - \|z\|^2}{2\sigma^2}\right) \\ &= \exp\left(-\frac{\|x\|^2}{2\sigma^2}\right) \exp\left(-\frac{\|z\|^2}{2\sigma^2}\right) \exp\left(\frac{x^T z}{\sigma^2}\right)\end{aligned}$$

- Use the Taylor expansion for  $\exp()$

$$\exp\left(\frac{x^T z}{\sigma^2}\right) = \sum_{n=0}^{\infty} \frac{(x^T z)^n}{\sigma^{2n} n!}$$

- Consider the Gaussian kernel

$$\begin{aligned}\exp\left(\frac{-\|x - z\|^2}{2\sigma^2}\right) &= \exp\left(\frac{-(x - z)^T(x - z)}{2\sigma^2}\right) \\ &= \exp\left(\frac{-\|x\|^2 + 2x^T z - \|z\|^2}{2\sigma^2}\right) \\ &= \exp\left(-\frac{\|x\|^2}{2\sigma^2}\right) \exp\left(-\frac{\|z\|^2}{2\sigma^2}\right) \exp\left(\frac{x^T z}{\sigma^2}\right)\end{aligned}$$

- Use the Taylor expansion for  $\exp()$

$$\exp\left(\frac{x^T z}{\sigma^2}\right) = \sum_{n=0}^{\infty} \frac{(x^T z)^n}{\sigma^{2n} n!}$$

Polynomial kernels of every degree!

- Bigger feature space increases the possibility of overfitting
  - Large margin solutions may still generalize reasonably well
- Alternative: add “penalties” to the objective to disincentivize complicated solutions

- Allow misclassification
  - Penalize misclassification linearly (just like in the perceptron algorithm)
    - Again, easier to work with than counting misclassifications
    - Objective stays convex
- Will let us handle data that isn't linearly separable!
- Idea: Take the constraints into the main objective
  - The objective function then becomes exactly like what we have seen in Perceptron/Linear Regression

# SVMs with Slack




$$\min_{w,b,\xi} \frac{1}{2} \|w\|^2 + c \sum_i \xi_i$$

such that

$$y_i(w^T x^{(i)} + b) \geq 1 - \xi_i, \text{ for all } i$$

$$\xi_i \geq 0, \text{ for all } i$$



$$\min_{w,b,\xi} \frac{1}{2} \|w\|^2 + c \sum_i \xi_i$$

such that

$$y_i(w^T x^{(i)} + b) \geq 1 - \xi_i, \text{ for all } i$$
$$\xi_i \geq 0, \text{ for all } i$$

Potentially allows some points to be misclassified/inside the margin



$$\min_{w,b,\xi} \frac{1}{2} \|w\|^2 + c \sum_i \xi_i$$

such that

$$y_i(w^T x^{(i)} + b) \geq 1 - \xi_i, \text{ for all } i$$

$$\xi_i \geq 0, \text{ for all } i$$

Constant  $c$  determines  
degree to which slack is  
penalized

$$\min_{w,b,\xi} \frac{1}{2} \|w\|^2 + c \sum_i \xi_i$$

such that

$$y_i(w^T x^{(i)} + b) \geq 1 - \xi_i, \text{ for all } i$$

$$\xi_i \geq 0, \text{ for all } i$$

- How does this objective change with  $c$ ?

$$\min_{w,b,\xi} \frac{1}{2} \|w\|^2 + c \sum_i \xi_i$$

such that

$$y_i(w^T x^{(i)} + b) \geq 1 - \xi_i, \text{ for all } i$$

$$\xi_i \geq 0, \text{ for all } i$$

- How does this objective change with  $c$ ?
  - As  $c \rightarrow \infty$ , requires a perfect classifier
  - As  $c \rightarrow 0$ , allows arbitrary classifiers (i.e., ignores the data)

$$\min_{w,b,\xi} \frac{1}{2} \|w\|^2 + c \sum_i \xi_i$$

such that

$$y_i(w^T x^{(i)} + b) \geq 1 - \xi_i, \text{ for all } i$$

$$\xi_i \geq 0, \text{ for all } i$$

- How should we pick  $c$ ?

$$\min_{w,b,\xi} \frac{1}{2} \|w\|^2 + c \sum_i \xi_i$$

such that

$$y_i(w^T x^{(i)} + b) \geq 1 - \xi_i, \text{ for all } i$$

$$\xi_i \geq 0, \text{ for all } i$$

- How should we pick  $c$ ?
  - Divide the data into three pieces training, testing, and **validation**
  - Use the validation set to tune the value of the **hyperparameter**  $c$

- General learning strategy
  - Build a classifier using the training data  $\leftarrow$  Model Parameters
  - Select hyperparameters using validation data  $\leftarrow$  C
  - Evaluate the chosen model with the selected hyperparameters on the test data

How can we tell if we overfit the training data?

- Gather Data + Labels
- Select feature vectors
- Randomly split into three groups
  - Training set
  - Validation set
  - Test set
- Experimentation cycle
  - Select a “good” hypothesis from the hypothesis space
  - Tune hyper-parameters using validation set
  - Compute accuracy on test set (fraction of correctly classified instances)

80%  
10%  
10%

$$\min_{w,b,\xi} \frac{1}{2} \|w\|^2 + c \sum_i \xi_i$$

such that

$$y_i(w^T x^{(i)} + b) \geq 1 - \xi_i, \text{ for all } i$$

$$\xi_i \geq 0, \text{ for all } i$$

- What is the optimal value of  $\xi$  for fixed  $w$  and  $b$ ?




$$\min_{w,b,\xi} \frac{1}{2} \|w\|^2 + c \sum_i \xi_i$$

such that

$$y_i(w^T x^{(i)} + b) \geq 1 - \xi_i, \text{ for all } i$$

$$\xi_i \geq 0, \text{ for all } i$$



- What is the optimal value of  $\xi$  for fixed  $w$  and  $b$ ?

- If  $y_i(w^T x^{(i)} + b) \geq 1$ , then  $\xi_i = 0$

- If  $y_i(w^T x^{(i)} + b) < 1$ , then  $\xi_i = 1 - y_i(w^T x^{(i)} + b)$

$\xi \rightarrow$  closed form sol?

$$\min_{w,b,\xi} \frac{1}{2} \|w\|^2 + c \sum_i \xi_i$$

such that

$$y_i(w^T x^{(i)} + b) \geq 1 - \xi_i, \text{ for all } i$$

$$\xi_i \geq 0, \text{ for all } i$$

- We can formulate this slightly differently
  - $\xi_i = \max\{0, 1 - y_i(w^T x^{(i)} + b)\}$
  - Does this look familiar?
  - Hinge loss provides an upper bound on Hamming loss

# Hinge Loss Formulation



- Obtain a new objective by substituting in for  $\xi$

$$\min_{w,b} \frac{1}{2} \|w\|^2 + c \sum_i \max\{0, 1 - y_i(w^T x^{(i)} + b)\}$$

Can minimize with gradient descent!

# Hinge Loss Formulation



- Obtain a new objective by substituting in for  $\xi$

$$\min_{w,b} \underbrace{\frac{1}{2} \|w\|^2}_{\text{Penalty to prevent overfitting}} + c \underbrace{\sum_i \max\{0, 1 - y_i(w^T x^{(i)} + b)\}}_{\text{Hinge loss}}$$

Penalty to prevent  
overfitting

"Avoids over fitting to  
the Data"

Hinge loss

"Fit well to the Data"

# REGULARIZATION!



- Until now, we have seen the following optimization problems:

$$\min_{w,b} \sum_i L(f(x^{(i)}, w, b), y_i) \leftarrow \text{Fit well to the Train Data.}$$


- In the case of Linear regression,  $L$  was the squared loss
- In Perceptron,  $L$  was Perceptron Loss
- The regularized version of this is:

$$\min_{w,b} \frac{1}{2} \|w\|^2 + c \sum_i L(f(x^{(i)}, w, b), y_i)$$


- $c$  is a hyper-parameter (again, to be tuned on validation set)

Why does  $\|w\|^2$  prevent overfitting?

$\|w\| \gg \gg \|w\|$



# Perceptron vs Hinge vs Square vs Zero-One Loss




- If the data is imbalanced (i.e., more positive examples than negative examples), may want to evenly distribute the error between the two classes

$$\min_{w,b,\xi} \frac{1}{2} \|w\|^2 + \frac{c}{N_+} \sum_{i:y_i=1} \xi_i + \frac{c}{N_-} \sum_{i:y_i=-1} \xi_i$$

such that

$$y_i(w^T x^{(i)} + b) \geq 1 - \xi_i, \text{ for all } i$$


$$\xi_i \geq 0, \text{ for all } i$$





- We argued, intuitively, that SVMs generalize better than the perceptron algorithm
  - How can we make this precise?

Margin  $\rightarrow$  Generalization  
(Test set performance)

- Where are we headed?
  - Non-Parametric Methods
    - $k$  nearest neighbor
    - Decision trees
  - Probabilistic Methods 
    - Bayesian Methods
    - Naïve Bayes
    - Logistic Regression
  - Unsupervised Learning
    - Clustering