



CS 4375

Nearest Neighbor Methods

Rishabh Iyer
University of Texas at Dallas


Nearest Neighbor Methods




- Learning
 - Store all training examples
- Classifying a new point x'
 - Find the training example $(x^{(i)}, y^{(i)})$ such that $x^{(i)}$ is closest (for some notion of close) to x'
 - Classify x' with the label $y^{(i)}$

Linear Models (Perceptrons, linear SVMs)
 $y_{\text{pred}} = \text{Sign}(w^T x' + b) \rightarrow \text{Binary Classification}$


Nearest Neighbor Methods





Nearest Neighbor Methods




Nearest Neighbor Methods



Nearest Neighbor Methods




Nearest Neighbor Methods




k -nearest neighbor methods look at the k closest points in the training set and take a majority vote
(should choose k to be odd)

Nearest Neighbor Methods




k -nearest neighbor methods look at the k closest points in the training set and take a majority vote
(should choose k to be odd)

1-NN Example



20-NN Example



Nearest Neighbor Methods



- Applies to data sets with points in \mathbb{R}^d
 - Best for large data sets with only a few (< 20) attributes
- Advantages
 - Learning is easy
 - Can learn complicated decision boundaries
- Disadvantages
 - Classification is slow (need to keep the entire training set around)
 - Easily fooled by irrelevant attributes

Practical Challenges



- How to choose the right measure of closeness?
 - Euclidean distance is popular, but many other possibilities
- How to pick k ?
 - Too small and the estimates are noisy, too large and the accuracy suffers
- What if the nearest neighbor is really far away?

$$\begin{array}{l} \begin{bmatrix} x_1^i \\ x_1 \end{bmatrix} \quad \begin{bmatrix} x_2^i \\ x_2 \end{bmatrix} \end{array} \rightarrow \begin{array}{l} \sqrt{\sum_{i=1}^d [x_1^i - x_2^i]^2} \rightarrow \text{L2 Distance} \\ \sum_{i=1}^d |x_1^i - x_2^i| \rightarrow \text{L1 Distance} \end{array}$$

$\|x_1 - x_2\|_2$
 $\|x_1 - x_2\|_1$

Choosing the Distance



- Euclidean distance makes sense when each of the features is roughly on the same scale
 - If the features are very different (e.g., height and age), then Euclidean distance makes less sense as height would be less significant than age simply because age has a larger range of possible values
- To correct for this, feature vectors are often recentered around their means and scaled by the standard deviation over the training set

Normalization



- Sample mean

$$\bar{x} = \frac{1}{n} \sum_{i=1}^n x^{(i)}$$

$$x_k^{\text{norm}} = \frac{x_k - \bar{x}_k}{\hat{\sigma}_k}$$

- Sample variance (biased)


$$E[x_k^{\text{norm}}] = E[x_k] - \bar{x}_k = 0$$

$$\hat{\sigma}_k^2 = \frac{1}{n} \sum_{i=1}^n (x_k^{(i)} - \bar{x}_k)^2$$

Irrelevant Attributes




Consider the nearest neighbor problem in one dimension





Irrelevant Attributes



Now, add a new attribute that is just random noise...





K-Dimensional Trees




- In order to do classification, we can compute the distances between all points in the training set and the point we are trying to classify
- With m data points in n -dimensional space, this takes $O(mn)$ time for Euclidean distance
$$\leftarrow \underset{\substack{\text{\# Feat}}}{O(n)} \times \underset{\substack{\text{\# points}}}{O(m)} \rightarrow \text{Single test Ex.}$$
- It is possible to do better if we do some preprocessing on the training data

- k-d trees provide a data structure that can help simplify the classification task by constructing a tree that partitions the search space
 - Starting with the entire training set, choose some dimension, i
 - Select an element of the training data whose i^{th} dimension has the median value among all elements of the training set
 - Divide the training set into two pieces: depending on whether their i^{th} attribute is smaller or larger than the median
 - Repeat this partitioning process on each of the two new pieces separately

K-Dimensional Trees





- Building a K-D tree from training data:
 - $\{(1,9), (2,3), (4,1), (3,7), (5,4), (6,8), (7,2), (8,8), (7,9), (9,6)\}$
 - pick random dimension, find median, split data, repeat



K-Dimensional Trees





- Building a K-D tree from training data:
 - $\{(1,9), (2,3), (4,1), (3,7), (5,4), (6,8), (7,2), (8,8), (7,9), (9,6)\}$
 - pick random dimension, find median, split data, repeat



K-Dimensional Trees: Inference



- Find NNs for new point (7,4)
 - find region containing (7,4)
 - compare to all points in region



- By design, the constructed k-d tree is “bushy”
 - The idea is that if new points to classify are evenly distributed throughout the space, then the expected (amortized) cost of classification is approximately $O(d \log n)$ operations
- Summary
 - k-NN is fast and easy to implement
 - No training required
 - Can be good in practice (where applicable)

\uparrow \uparrow
Dim # Train Examples