



Practical ML Advice

Rishabh Iyer

Proper Experimental Methodology Can Have a Huge Impact:

A 2002 paper in *Nature* (a major journal) needed to be corrected due to “training on the testing set”

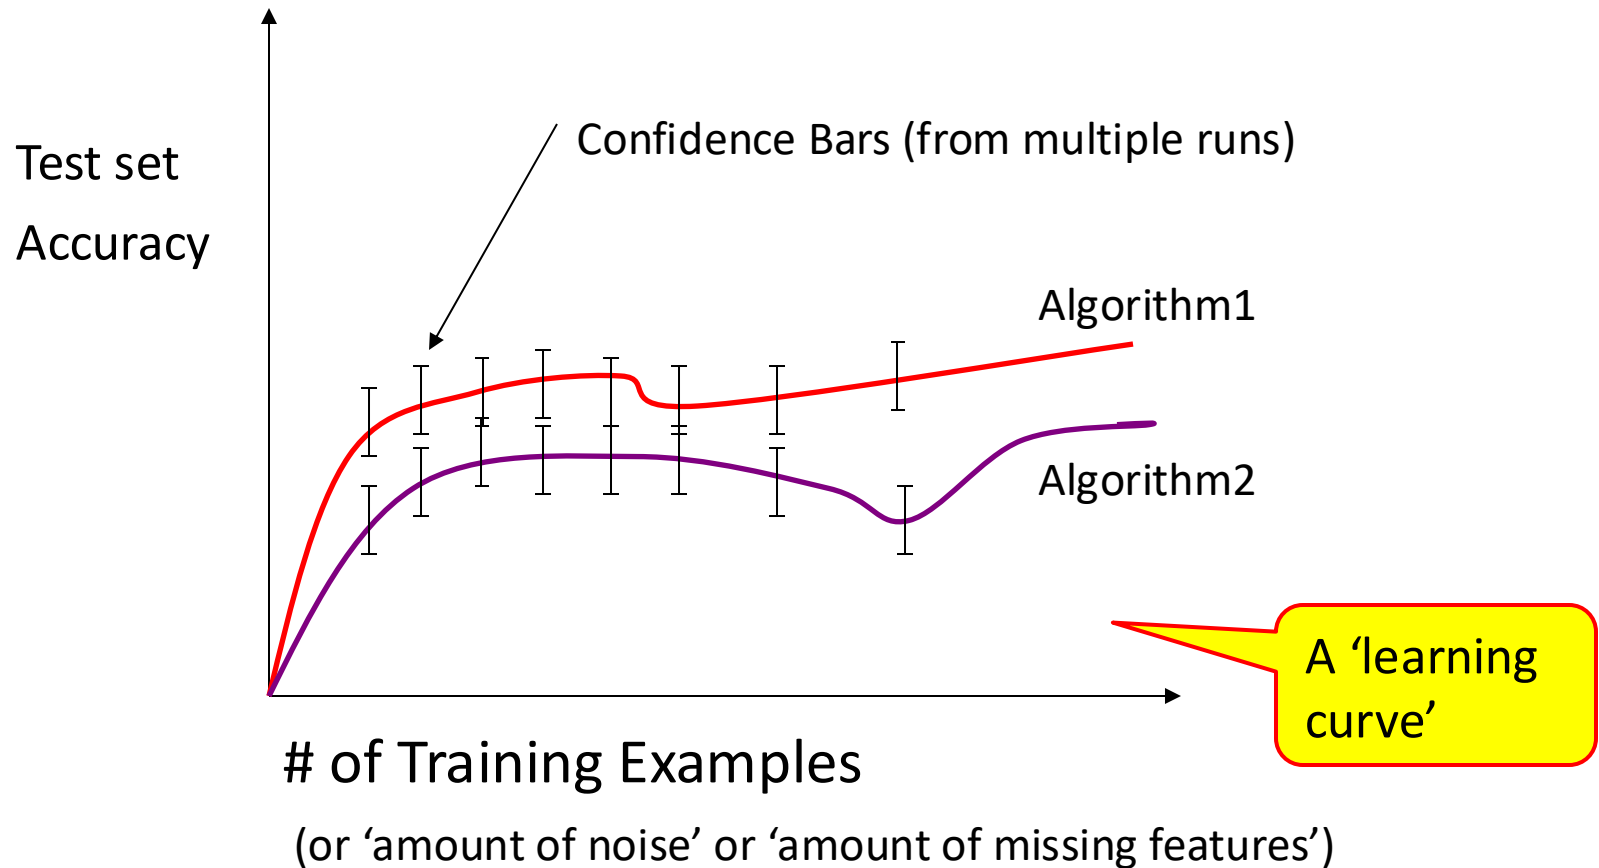
Original report : 95% accuracy (5% error rate)

Corrected report (which still is buggy):

73% accuracy (27% error rate)

Error rate increased over 400%!!!

Some Typical ML Experiments



Typical Experiments



| | Test Set Performance |
|------------------|----------------------|
| Full System | 80% |
| Without Module A | 75% |
| Without Module B | 62% |
| | |

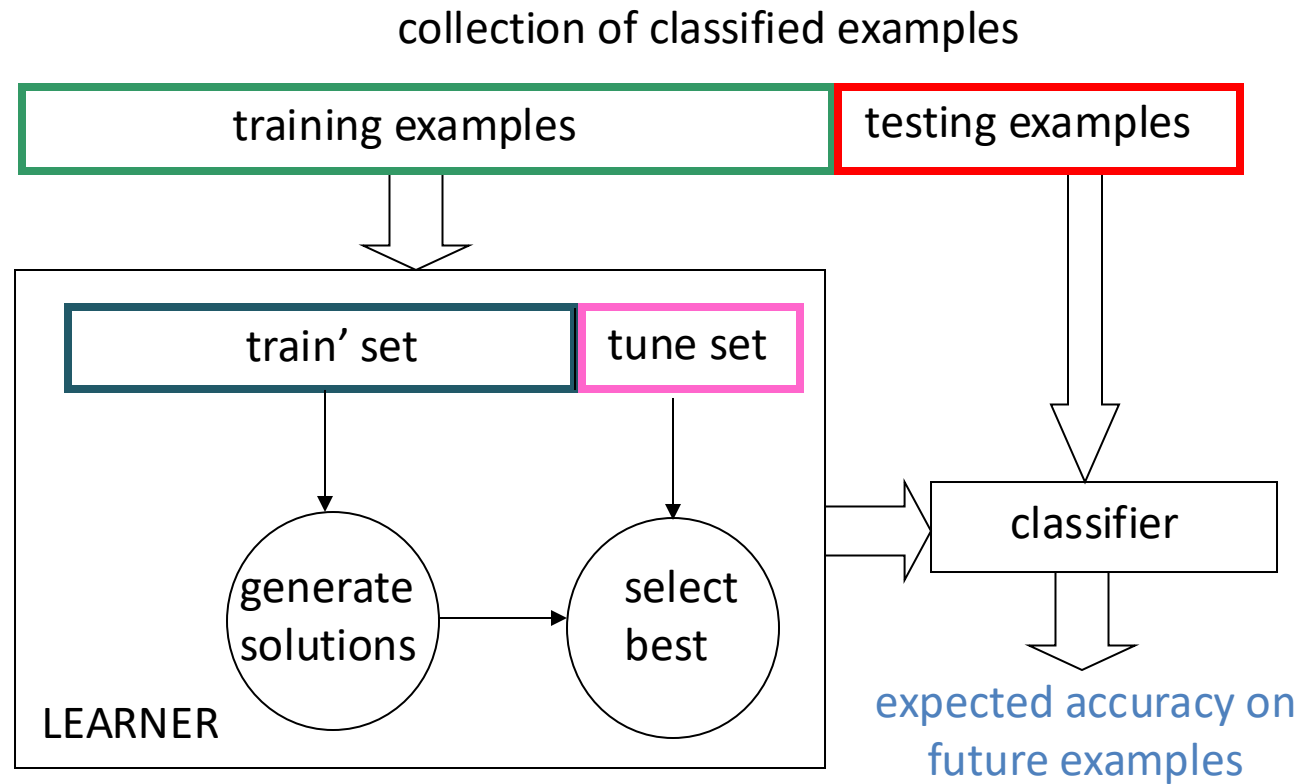
- 1) Start with a dataset of labeled examples
- 2) Randomly partition into N groups
- 3a) N times, combine $N - 1$ groups into a train set
- 3b) Provide **training set** to learning system
- 3c) Measure accuracy on “left out” group (the **test set**)



Called **N -fold cross validation**

- Often, an ML system has to choose when to stop learning, select among alternative answers, etc.
- One wants the model that produces the highest accuracy on **future** examples (“overfitting avoidance”)
- It is a “**cheat**” to look at the **test** set while still learning
- Better method
 - Set aside part of the training set
 - Measure performance on this validation data to estimate future performance for a given set of hyperparameters
 - Use best hyperparameter settings, train with **all** training data (except **test** set) to estimate future performance on **new** examples

A typical Learning system



Statistical techniques such as 10-fold cross validation and *t*-tests are used to get meaningful results

Multiple Tuning sets



- Using a **single** tuning set can be unreliable predictor, plus some data “wasted”
 - 1) For each possible set of hyperparameters
 - a) Divide training data into **train** and **valid.** sets, using ***N*-fold cross validation**
 - b) Score this set of hyperparameter values: average **valid.** set accuracy over the *N* folds
 - 2) Use **best** set of hyperparameter settings and **all** (train + valid.) examples
 - 3) Apply resulting model to **test** set



EVALUATING ML MODELS

Regression Evaluation Metrics



1. Mean Absolute Error (MAE)

$$\text{MAE} = \frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i|$$

- Measures the average magnitude of errors in predictions.
- **Robust to outliers**, all errors are treated equally.

2. Mean Squared Error (MSE)

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

- Squares the errors before averaging.
- **Penalizes larger errors** more than smaller ones.
- Sensitive to outliers.

3. Root Mean Squared Error (RMSE)

$$\text{RMSE} = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2}$$

- Square root of MSE.
- **Interpreted in the same units** as the target variable.
- Useful when **larger errors are more significant**.

4. R-squared (R^2 Score)

$$R^2 = 1 - \frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{\sum_{i=1}^n (y_i - \bar{y})^2}$$

- Measures the **proportion of variance explained** by the model.
- Ranges from $-\infty$ to 1:
 - $R^2 = 1$: Perfect prediction
 - $R^2 = 0$: Model no better than the mean
 - $R^2 < 0$: Worse than predicting the mean

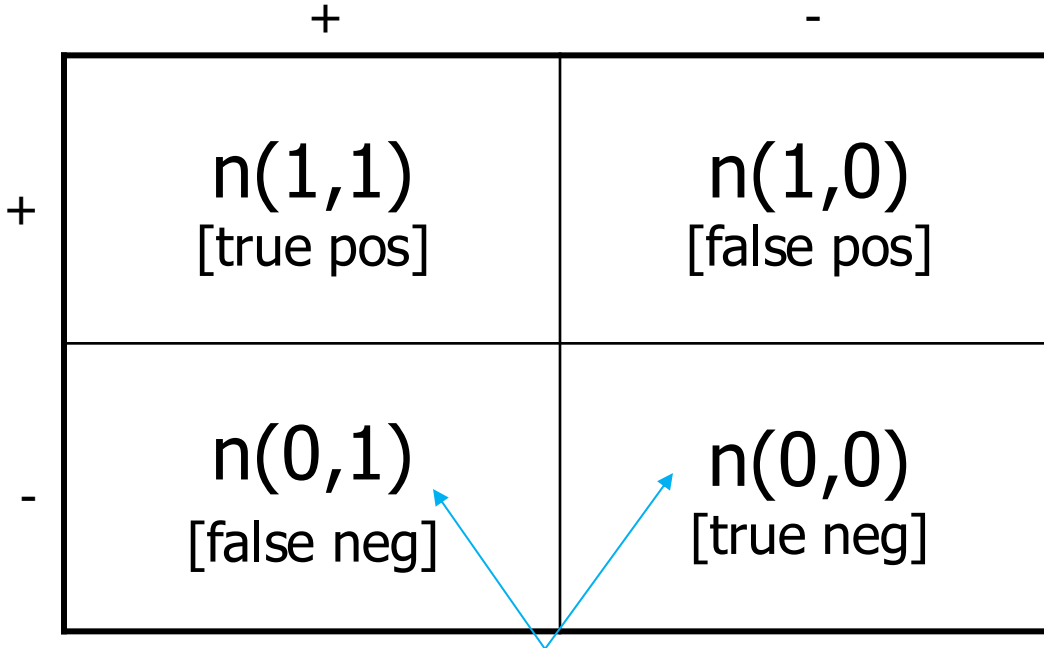
Contingency Tables



(special case of 'confusion matrices')

| | | True Answer | |
|-------------------------|---|-------------------------|-------------------------|
| | | + | - |
| Algorithm Answer | + | $n(1,1)$ [true pos] | $n(1,0)$ [false pos] |
| | - | $n(0,1)$ [false neg] | $n(0,0)$ [true neg] |

Counts of occurrences



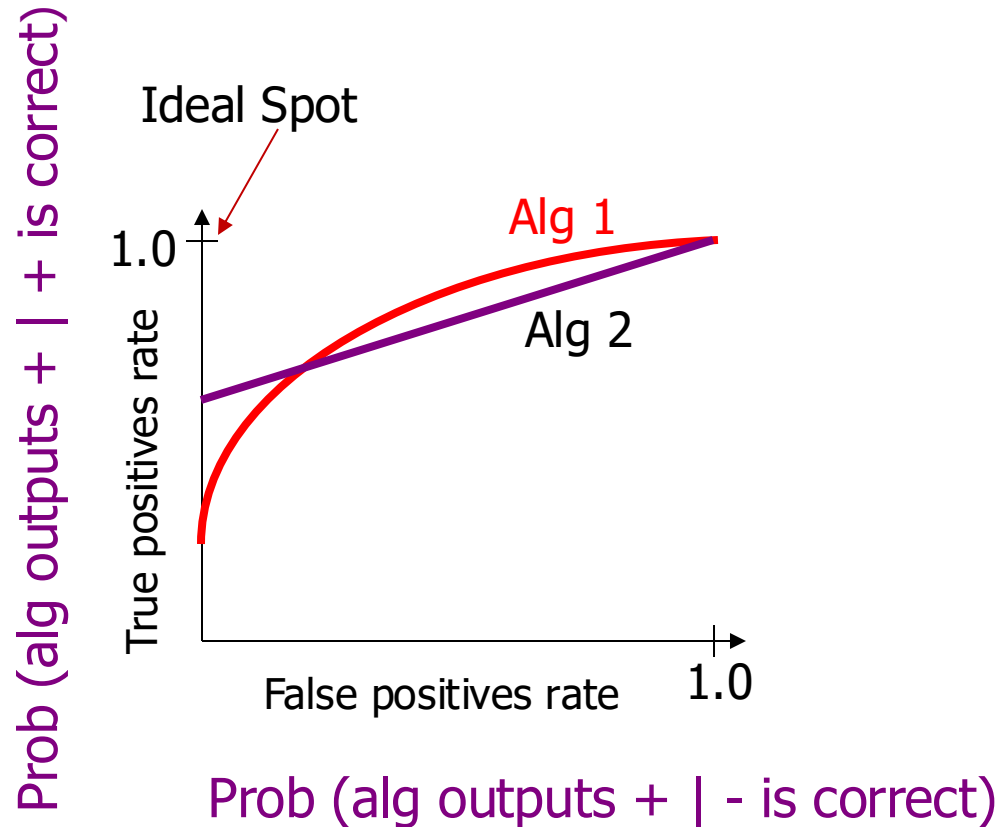
True Positive Rate
(TPR) $= n(1,1) / (n(1,1) + n(0,1))$
 $=$ correctly categorized +’s / total positives
 $\sim P(\text{algo outputs } + \mid + \text{ is correct})$

False Positive Rate
(FPR) $= n(1,0) / (n(1,0) + n(0,0))$
 $=$ incorrectly categorized –’s / total neg’s
 $\sim P(\text{algo outputs } + \mid - \text{ is correct})$

Can similarly define False Negative Rate and True Negative Rate

- ROC: *Receiver Operating Characteristics*
- Started for radar research during WWII
- Judging algorithms on accuracy alone may not be good enough when **getting a positive wrong** costs more than **getting a negative wrong** (or vice versa)
 - e.g., medical tests for serious diseases
 - e.g., a movie-recommender system

ROC Curves Graphically



Different algorithms can work better in different parts of ROC space. This depends on cost of false + vs false -

Creating an ROC Curve

The Standard Approach:

- You need an ML algorithm that outputs **NUMERIC** results such as `prob(example is +)`
- You can use ensemble methods to get this from a model that only provides Boolean outputs
 - e.g., have 100 models vote & count votes

Alg. for Creating ROC Curves

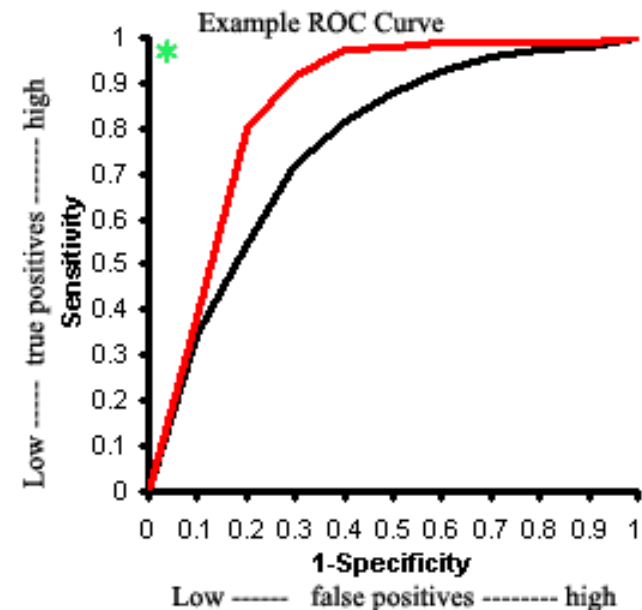


Step 1: Sort predictions on test set

Step 2: Locate a *threshold* between examples with opposite categories

Step 3: Compute TPR & FPR for each threshold of Step 2

Step 4: Connect the dots

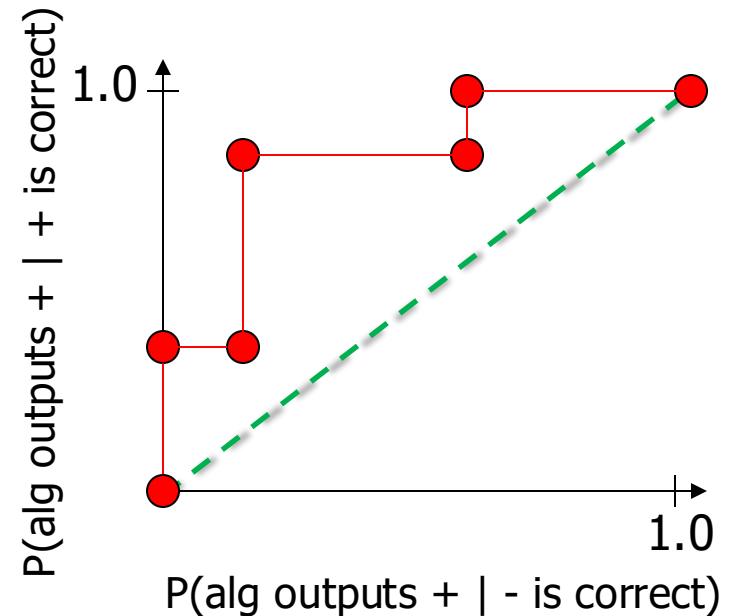


Plotting ROC Curves - Example



ML Algo Output (Sorted) Correct Category

| | | | |
|-------|-----|----------------------|---|
| Ex 9 | .99 | | + |
| Ex 7 | .98 | TPR=(2/5), FPR=(0/5) | + |
| Ex 1 | .72 | TPR=(2/5), FPR=(1/5) | - |
| Ex 2 | .70 | | + |
| Ex 6 | .65 | TPR=(4/5), FPR=(1/5) | + |
| Ex 10 | .51 | | - |
| Ex 3 | .39 | TPR=(4/5), FPR=(3/5) | - |
| Ex 5 | .24 | TPR=(5/5), FPR=(3/5) | + |
| Ex 4 | .11 | | - |
| Ex 8 | .01 | TPR=(5/5), FPR=(5/5) | - |

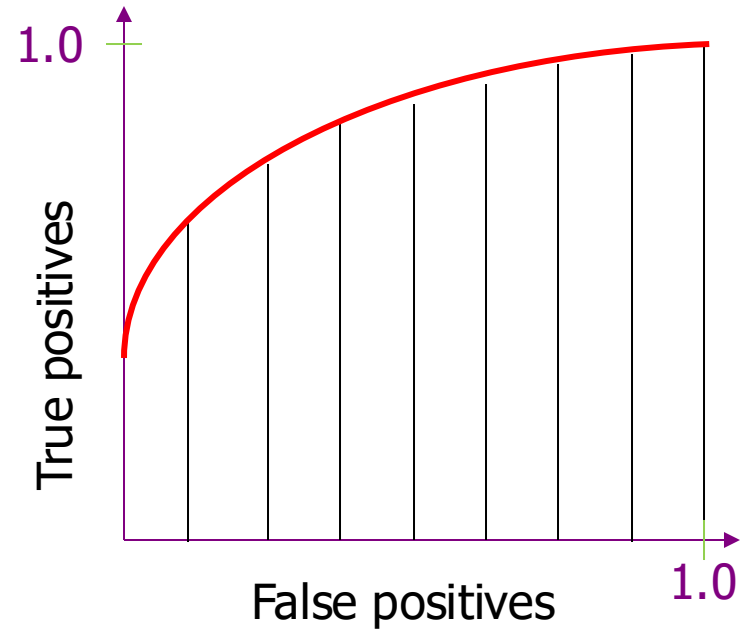


Algorithm predicts + if its output is ≥ 0

Area Under ROC Curve



- A common metric for experiments is to numerically integrate the ROC Curve
 - Usually called AUC
 - Probability that ML alg. will “rank” a randomly chosen positive instance higher than a randomly chosen negative one
 - Can summarize the curve **too much** in practice



Asymmetric Error Costs



- Assume that $\text{cost}(FP) \neq \text{cost}(FN)$
- You would like to pick a threshold that minimizes

$$\begin{aligned} E(\text{total cost}) \\ = \text{cost}(FP) \times \text{pr}(FP) \times (\# \text{ of neg ex's}) + \\ \text{cost}(FN) \times \text{pr}(FN) \times (\# \text{ of pos ex's}) \end{aligned}$$

- You could also have (maybe negative) costs for TP and TN (assumed zero in above)

- One strength of ROC curves is that they are a good way to deal with **skewed** data ($|+| \gg |-|$) since the axes are fractions (rates) independent of the # of examples
- You must be careful though!
 - Low FPR * (many negative ex) = **sizable number of FP**
 - Possibly more than # of TP

Precision vs. Recall

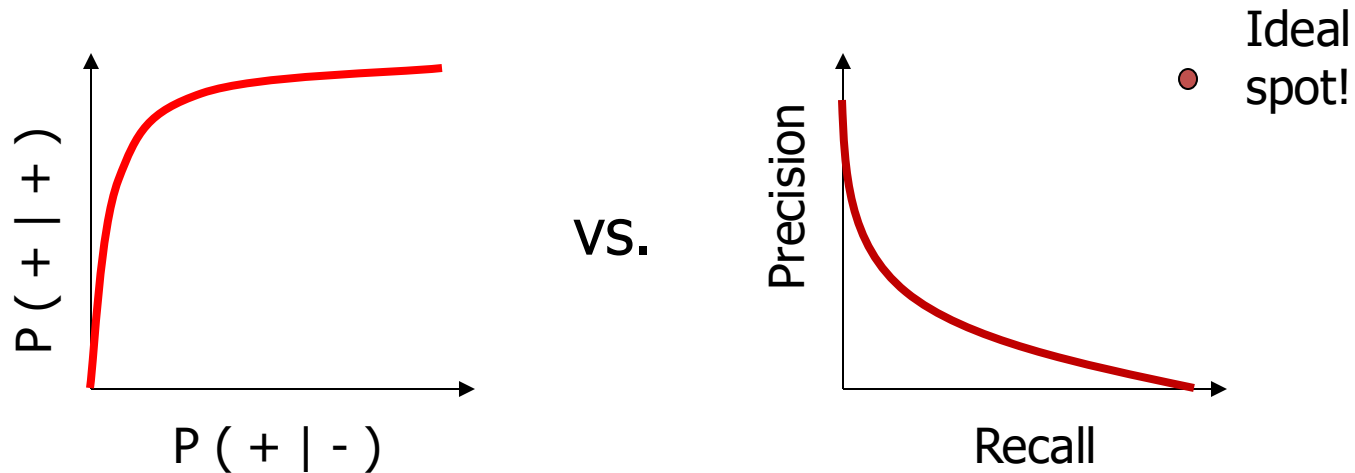


- Think about search engines...
- **Precision** = (# of relevant items retrieved)
/ (total # of items retrieved)
 $= n(1,1) / (n(1,1) + n(1,0))$
 $\cong P(\text{is pos} \mid \text{called pos})$
- **Recall** = (# of relevant items retrieved)
/ (# of relevant items that exist)
 $= n(1,1) / (n(1,1) + n(0,1)) = \text{TPR}$
 $\cong P(\text{called pos} \mid \text{is pos})$
- Notice that $n(0,0)$ is not used in either formula
Therefore you get no credit for filtering out irrelevant items

ROC vs. Precision-Recall



You can get very different visual results
on the same data!



The reason for this is that there may be lots of – ex's
(e.g., might need to include 100 neg's to get 1 more pos)

The F1 Measure



- Figure of merit that combines precision and recall

$$F_1 = 2 \cdot \frac{P \cdot R}{P + R}$$

where P = precision; R = recall. This is twice the harmonic mean of P and R .

- We can plot $F1$ as a function of the classification threshold θ

Summary of Evaluation Metrics



| Metric | Best Use Case |
|----------------------|--|
| Accuracy | Balanced classes with equal error cost |
| F1 Score | Imbalanced classes, equal importance to precision/recall |
| Precision | When false positives are costly (e.g., spam filtering) |
| Recall | When false negatives are costly (e.g., medical diagnosis) |
| AUC | Imbalanced data, ranking-focused applications |
| MAE (L1) | Regression with outliers , stable performance |
| RMSE (L2) | Regression where larger errors are worse |
| R ² Score | Overall model fit in regression |